

# lp5

May 6, 2025

EXP - 6

Linear regression by using Deep Neural network: Implement Boston housing price prediction problem by Linear regression using Deep Neural network. Use Boston House price prediction dataset

```
[11]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import pandas as pd
```

```
[16]: # import data

data = pd.read_csv(r'/Users/sakshisachinpotdar/Downloads/1_boston_housing.csv')
print(data.head())
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	

	b	lstat	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

```
[18]: # SPECIFY X & Y
# then train_test

y = data['MEDV']
x = data.drop(columns = ['MEDV'])
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2,
↪random_state = 32)
```

```
[21]: # standardise fts
```

```
scaler = StandardScaler()  
x_train = scaler.fit_transform(x_train)  
x_test = scaler.transform(x_test)
```

```
[25]: # BUILD MODEL
```

```
model = Sequential()  
  
model.add(Dense(64, input_dim = x_train.shape[1], activation = 'relu')) # 1st_  
    ↳hidden layer with 64 neurons  
model.add(Dense(32, activation = 'relu' )) # 2nd hidden layer with 32 neurons  
model.add(Dense(1)) #Single output
```

```
[26]: # MODEL COMPILATION
```

```
model.compile(optimizer = 'adam', loss = 'mean_squared_error')
```

```
[27]: # TRAIN MODEL
```

```
model.fit(x_train, y_train, epochs = 100, batch_size = 32, validation_data =_  
    ↳(x_test, y_test))
```

```
Epoch 1/100  
13/13          0s 5ms/step - loss:  
574.0211 - val_loss: 530.7908  
Epoch 2/100  
13/13          0s 2ms/step - loss:  
502.6110 - val_loss: 466.1854  
Epoch 3/100  
13/13          0s 2ms/step - loss:  
395.8426 - val_loss: 386.9705  
Epoch 4/100  
13/13          0s 2ms/step - loss:  
342.1078 - val_loss: 293.5982  
Epoch 5/100  
13/13          0s 2ms/step - loss:  
239.8614 - val_loss: 204.0109  
Epoch 6/100  
13/13          0s 3ms/step - loss:  
175.3199 - val_loss: 141.3174  
Epoch 7/100  
13/13          0s 2ms/step - loss:  
98.3913 - val_loss: 110.9292  
Epoch 8/100  
13/13          0s 2ms/step - loss:  
78.4389 - val_loss: 90.1927
```

Epoch 9/100  
13/13 0s 2ms/step - loss:  
57.6758 - val\_loss: 72.1971  
Epoch 10/100  
13/13 0s 2ms/step - loss:  
45.9713 - val\_loss: 59.1600  
Epoch 11/100  
13/13 0s 2ms/step - loss:  
34.1701 - val\_loss: 50.5682  
Epoch 12/100  
13/13 0s 2ms/step - loss:  
30.2187 - val\_loss: 44.5006  
Epoch 13/100  
13/13 0s 2ms/step - loss:  
26.8668 - val\_loss: 40.2278  
Epoch 14/100  
13/13 0s 2ms/step - loss:  
32.7378 - val\_loss: 36.9662  
Epoch 15/100  
13/13 0s 2ms/step - loss:  
23.5899 - val\_loss: 34.6057  
Epoch 16/100  
13/13 0s 2ms/step - loss:  
23.8061 - val\_loss: 32.1485  
Epoch 17/100  
13/13 0s 2ms/step - loss:  
23.0260 - val\_loss: 30.5701  
Epoch 18/100  
13/13 0s 2ms/step - loss:  
22.0493 - val\_loss: 29.0889  
Epoch 19/100  
13/13 0s 2ms/step - loss:  
21.7630 - val\_loss: 27.9306  
Epoch 20/100  
13/13 0s 2ms/step - loss:  
20.3039 - val\_loss: 26.7396  
Epoch 21/100  
13/13 0s 2ms/step - loss:  
15.8878 - val\_loss: 25.7958  
Epoch 22/100  
13/13 0s 2ms/step - loss:  
16.0272 - val\_loss: 25.1617  
Epoch 23/100  
13/13 0s 2ms/step - loss:  
19.3797 - val\_loss: 24.3228  
Epoch 24/100  
13/13 0s 2ms/step - loss:  
19.1960 - val\_loss: 23.8775

Epoch 25/100  
13/13 0s 2ms/step - loss:  
16.1112 - val\_loss: 23.0935  
Epoch 26/100  
13/13 0s 2ms/step - loss:  
14.5181 - val\_loss: 22.7338  
Epoch 27/100  
13/13 0s 2ms/step - loss:  
16.4957 - val\_loss: 21.9592  
Epoch 28/100  
13/13 0s 2ms/step - loss:  
13.3504 - val\_loss: 21.9457  
Epoch 29/100  
13/13 0s 2ms/step - loss:  
13.3903 - val\_loss: 21.3593  
Epoch 30/100  
13/13 0s 2ms/step - loss:  
12.0795 - val\_loss: 21.2926  
Epoch 31/100  
13/13 0s 4ms/step - loss:  
13.2743 - val\_loss: 20.7340  
Epoch 32/100  
13/13 0s 2ms/step - loss:  
13.1588 - val\_loss: 20.6477  
Epoch 33/100  
13/13 0s 2ms/step - loss:  
11.8908 - val\_loss: 20.3710  
Epoch 34/100  
13/13 0s 2ms/step - loss:  
13.2728 - val\_loss: 20.0850  
Epoch 35/100  
13/13 0s 2ms/step - loss:  
13.6316 - val\_loss: 20.0262  
Epoch 36/100  
13/13 0s 2ms/step - loss:  
16.2454 - val\_loss: 19.7193  
Epoch 37/100  
13/13 0s 2ms/step - loss:  
12.4100 - val\_loss: 19.5555  
Epoch 38/100  
13/13 0s 2ms/step - loss:  
12.8917 - val\_loss: 19.6123  
Epoch 39/100  
13/13 0s 2ms/step - loss:  
13.6737 - val\_loss: 19.4560  
Epoch 40/100  
13/13 0s 2ms/step - loss:  
11.5287 - val\_loss: 19.3714

Epoch 41/100  
13/13 0s 2ms/step - loss:  
13.4470 - val\_loss: 18.9986  
Epoch 42/100  
13/13 0s 2ms/step - loss:  
9.5816 - val\_loss: 19.1964  
Epoch 43/100  
13/13 0s 2ms/step - loss:  
9.8440 - val\_loss: 18.9520  
Epoch 44/100  
13/13 0s 4ms/step - loss:  
9.7949 - val\_loss: 18.9620  
Epoch 45/100  
13/13 0s 2ms/step - loss:  
12.3425 - val\_loss: 18.8336  
Epoch 46/100  
13/13 0s 2ms/step - loss:  
10.8069 - val\_loss: 18.8129  
Epoch 47/100  
13/13 0s 2ms/step - loss:  
10.6927 - val\_loss: 18.6624  
Epoch 48/100  
13/13 0s 2ms/step - loss:  
9.7675 - val\_loss: 18.6233  
Epoch 49/100  
13/13 0s 2ms/step - loss:  
11.4991 - val\_loss: 18.6640  
Epoch 50/100  
13/13 0s 2ms/step - loss:  
9.4074 - val\_loss: 18.5716  
Epoch 51/100  
13/13 0s 2ms/step - loss:  
8.8536 - val\_loss: 18.5260  
Epoch 52/100  
13/13 0s 2ms/step - loss:  
9.8002 - val\_loss: 18.4020  
Epoch 53/100  
13/13 0s 2ms/step - loss:  
12.8017 - val\_loss: 18.3972  
Epoch 54/100  
13/13 0s 2ms/step - loss:  
12.3081 - val\_loss: 18.5454  
Epoch 55/100  
13/13 0s 2ms/step - loss:  
8.5845 - val\_loss: 18.3113  
Epoch 56/100  
13/13 0s 2ms/step - loss:  
10.1593 - val\_loss: 18.1546

Epoch 57/100  
13/13 0s 3ms/step - loss:  
9.5331 - val\_loss: 18.2164  
Epoch 58/100  
13/13 0s 2ms/step - loss:  
9.9207 - val\_loss: 18.2329  
Epoch 59/100  
13/13 0s 2ms/step - loss:  
9.7162 - val\_loss: 18.1948  
Epoch 60/100  
13/13 0s 2ms/step - loss:  
9.6262 - val\_loss: 18.2148  
Epoch 61/100  
13/13 0s 2ms/step - loss:  
10.6908 - val\_loss: 18.2444  
Epoch 62/100  
13/13 0s 2ms/step - loss:  
11.5055 - val\_loss: 18.0579  
Epoch 63/100  
13/13 0s 2ms/step - loss:  
9.5257 - val\_loss: 18.1005  
Epoch 64/100  
13/13 0s 3ms/step - loss:  
12.5958 - val\_loss: 18.0821  
Epoch 65/100  
13/13 0s 2ms/step - loss:  
8.7596 - val\_loss: 18.0166  
Epoch 66/100  
13/13 0s 2ms/step - loss:  
8.2777 - val\_loss: 18.0720  
Epoch 67/100  
13/13 0s 2ms/step - loss:  
9.3511 - val\_loss: 18.0093  
Epoch 68/100  
13/13 0s 2ms/step - loss:  
7.5161 - val\_loss: 17.9276  
Epoch 69/100  
13/13 0s 2ms/step - loss:  
7.9653 - val\_loss: 17.7819  
Epoch 70/100  
13/13 0s 2ms/step - loss:  
7.5801 - val\_loss: 17.7847  
Epoch 71/100  
13/13 0s 2ms/step - loss:  
10.7447 - val\_loss: 17.8071  
Epoch 72/100  
13/13 0s 2ms/step - loss:  
7.9200 - val\_loss: 17.8893

Epoch 73/100  
13/13 0s 2ms/step - loss:  
10.2547 - val\_loss: 17.9021  
Epoch 74/100  
13/13 0s 3ms/step - loss:  
8.1569 - val\_loss: 17.8021  
Epoch 75/100  
13/13 0s 2ms/step - loss:  
7.7583 - val\_loss: 17.7611  
Epoch 76/100  
13/13 0s 2ms/step - loss:  
10.1824 - val\_loss: 17.6027  
Epoch 77/100  
13/13 0s 2ms/step - loss:  
9.4185 - val\_loss: 17.7168  
Epoch 78/100  
13/13 0s 2ms/step - loss:  
9.4211 - val\_loss: 17.8408  
Epoch 79/100  
13/13 0s 2ms/step - loss:  
9.3143 - val\_loss: 17.6500  
Epoch 80/100  
13/13 0s 2ms/step - loss:  
8.2847 - val\_loss: 17.7216  
Epoch 81/100  
13/13 0s 2ms/step - loss:  
8.8211 - val\_loss: 17.7020  
Epoch 82/100  
13/13 0s 2ms/step - loss:  
8.5291 - val\_loss: 17.7535  
Epoch 83/100  
13/13 0s 3ms/step - loss:  
8.6597 - val\_loss: 17.6259  
Epoch 84/100  
13/13 0s 2ms/step - loss:  
7.6142 - val\_loss: 17.6598  
Epoch 85/100  
13/13 0s 2ms/step - loss:  
6.6553 - val\_loss: 17.8108  
Epoch 86/100  
13/13 0s 2ms/step - loss:  
7.4534 - val\_loss: 17.5466  
Epoch 87/100  
13/13 0s 2ms/step - loss:  
9.1237 - val\_loss: 17.6228  
Epoch 88/100  
13/13 0s 2ms/step - loss:  
7.3712 - val\_loss: 17.5148

```

Epoch 89/100
13/13          0s 2ms/step - loss:
7.7251 - val_loss: 17.6035
Epoch 90/100
13/13          0s 2ms/step - loss:
7.8541 - val_loss: 17.5903
Epoch 91/100
13/13          0s 3ms/step - loss:
7.8447 - val_loss: 17.4894
Epoch 92/100
13/13          0s 2ms/step - loss:
10.6556 - val_loss: 17.4689
Epoch 93/100
13/13          0s 2ms/step - loss:
8.1113 - val_loss: 17.5697
Epoch 94/100
13/13          0s 2ms/step - loss:
9.6750 - val_loss: 17.4576
Epoch 95/100
13/13          0s 2ms/step - loss:
8.1899 - val_loss: 17.6628
Epoch 96/100
13/13          0s 2ms/step - loss:
8.8456 - val_loss: 17.3911
Epoch 97/100
13/13          0s 2ms/step - loss:
8.1951 - val_loss: 17.5429
Epoch 98/100
13/13          0s 2ms/step - loss:
7.6914 - val_loss: 17.4342
Epoch 99/100
13/13          0s 3ms/step - loss:
7.9066 - val_loss: 17.3363
Epoch 100/100
13/13          0s 2ms/step - loss:
6.7606 - val_loss: 17.4322

```

[27]: <keras.src.callbacks.history.History at 0x13f726570>

```

[29]: # Evaluate model

loss = model.evaluate(x_test, y_test)
print(f'Model loss (MSE): {loss}')

```

```

4/4          0s 3ms/step - loss:
19.1693
Model loss (MSE): 17.432201385498047

```



```
[31]: # make predictions

y_pred = model.predict(x_test)
print("Predictions: ", y_pred[:5]) # Print the first 5 predictions
```

```
4/4          0s 6ms/step
Predictions: [[13.574309]
 [28.991796]
 [15.981    ]
 [20.510895]
 [24.87274  ]]
```

```
[ ]:
```

Experiment 7

```
[2]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Embedding, LSTM, SpatialDropout1D
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.utils import to_categorical
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
import pandas as pd

# Load your CSV dataset
```

```
[3]: df = pd.read_csv('/Users/sakshisachinpotdar/Downloads/IMDB Dataset - IMDB_
↳Dataset.csv')
```

```
[4]: # Let's assume your CSV has two columns: 'review' and 'sentiment'
reviews = df['review'].values
labels = df['sentiment'].values
```

```
[5]: # Convert string labels ('positive' and 'negative') to numerical labels (1 and
↳0)
label_encoder = LabelEncoder()
y = label_encoder.fit_transform(labels) # Converts 'positive' to 1 and
↳'negative' to 0
```

```
[6]: # Tokenize text
tokenizer = Tokenizer(num_words=10000)
tokenizer.fit_on_texts(reviews)
X = tokenizer.texts_to_sequences(reviews)
X = pad_sequences(X, padding='post', maxlen=100) # Ensure same length
```

```
[7]: # Split into train and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

```
[10]: model = Sequential()
model.add(Embedding(input_dim=10000, output_dim=128))
model.add(SpatialDropout1D(0.2)) # Dropout to prevent overfitting
model.add(LSTM(100, dropout=0.2, recurrent_dropout=0.2))
model.add(Dense(1, activation='sigmoid'))
```

```
[11]: model.compile(loss='binary_crossentropy', optimizer=Adam(learning_rate=0.001),
↳metrics=['accuracy'])
```

```
[12]: history = model.fit(X_train, y_train, epochs=5, batch_size=64,
↳validation_data=(X_test, y_test), verbose=2)
```

```
Epoch 1/5
625/625 - 65s - 104ms/step - accuracy: 0.7822 - loss: 0.4666 - val_accuracy:
0.8347 - val_loss: 0.3775
Epoch 2/5
625/625 - 64s - 102ms/step - accuracy: 0.8614 - loss: 0.3345 - val_accuracy:
0.8549 - val_loss: 0.3418
Epoch 3/5
625/625 - 66s - 106ms/step - accuracy: 0.8881 - loss: 0.2798 - val_accuracy:
0.8665 - val_loss: 0.3202
Epoch 4/5
625/625 - 65s - 104ms/step - accuracy: 0.9037 - loss: 0.2436 - val_accuracy:
0.8694 - val_loss: 0.3383
Epoch 5/5
625/625 - 65s - 104ms/step - accuracy: 0.9182 - loss: 0.2133 - val_accuracy:
0.8706 - val_loss: 0.3314
```

```
[13]: loss, accuracy = model.evaluate(X_test, y_test)
print(f"Test Accuracy: {accuracy}")
```

```
313/313          5s 16ms/step -
accuracy: 0.8744 - loss: 0.3240
Test Accuracy: 0.8705999851226807
```

```
[ ]:
```

## EXPERIMENT 8

Convolutional neural network (CNN) Use MNIST Fashion Dataset and create a classifier to classify fashion clothing into categories.

```
[19]: import numpy as np
import matplotlib.pyplot as plt
```

```

from tensorflow.keras.datasets import fashion_mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, MaxPooling2D, Conv2D

# Load the Fashion MNIST dataset
(train_x, train_y), (test_x, test_y) = fashion_mnist.load_data()

# Normalize the data (important for good performance)
train_x = train_x.astype('float32') / 255.0
test_x = test_x.astype('float32') / 255.0

# Reshape the data to include the channel dimension (grayscale image)
train_x = train_x.reshape(-1, 28, 28, 1)
test_x = test_x.reshape(-1, 28, 28, 1)

# Define the model
model = Sequential()

# Add a Conv2D layer with 64 filters and ReLU activation function
model.add(Conv2D(filters=64, kernel_size=(3, 3), activation='relu',
    ↪input_shape=(28, 28, 1)))

# Add a MaxPooling2D layer to downsample the feature map
model.add(MaxPooling2D(pool_size=(2, 2)))

# Flatten the output to feed into a fully connected layer
model.add(Flatten())

# Add a Dense layer with 128 neurons and ReLU activation
model.add(Dense(128, activation='relu'))

# Add the output layer with 10 neurons (one for each class) and softmax
    ↪activation
model.add(Dense(10, activation='softmax'))

# Print the model summary
model.summary()

# Compile the model with the Adam optimizer and sparse categorical
    ↪cross-entropy loss
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
    ↪metrics=['accuracy'])

# Train the model with a validation split of 0.2
model.fit(train_x, train_y, epochs=5, validation_split=0.2)

# Evaluate the model on the test data

```

```

loss, acc = model.evaluate(test_x, test_y)
print(f"Test loss: {loss}, Test accuracy: {acc}")

# Define the label names corresponding to the 10 classes
labels = ['t-shirt', 'trouser', 'pullover', 'dress', 'coat', 'sandal', 'shirt', 'sneaker', 'bag', 'ankle_boots']

# Make predictions on the first test image
predictions = model.predict(test_x[:1])

# Map the predicted label to the correct class
label = labels[np.argmax(predictions)]

# Plot the first test image
print("Predicted label:", label)
plt.imshow(test_x[:1][0], cmap='gray') # Use 'gray' to show the image in grayscale
plt.show()

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d_11 (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 64)	0
flatten_1 (Flatten)	(None, 10816)	0
dense (Dense)	(None, 128)	1,384,576
dense_1 (Dense)	(None, 10)	1,290

Total params: 1,386,506 (5.29 MB)

Trainable params: 1,386,506 (5.29 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/5

1500/1500

11s 7ms/step -


accuracy: 0.8120 - loss: 0.5336 - val\_accuracy: 0.8932 - val\_loss: 0.2919

Epoch 2/5

```

1500/1500          11s 8ms/step -
accuracy: 0.9040 - loss: 0.2668 - val_accuracy: 0.9057 - val_loss: 0.2605
Epoch 3/5
1500/1500          11s 8ms/step -
accuracy: 0.9191 - loss: 0.2192 - val_accuracy: 0.9082 - val_loss: 0.2455
Epoch 4/5
1500/1500          11s 7ms/step -
accuracy: 0.9350 - loss: 0.1760 - val_accuracy: 0.9154 - val_loss: 0.2403
Epoch 5/5
1500/1500          11s 7ms/step -
accuracy: 0.9478 - loss: 0.1440 - val_accuracy: 0.9106 - val_loss: 0.2651
313/313           0s 1ms/step -
accuracy: 0.9034 - loss: 0.2861
Test loss: 0.2759930491447449, Test accuracy: 0.9075000286102295
1/1              0s 21ms/step
Predicted label: ankle_boots

```



[ ]:

## EXPERIMENT 9

Recurrent neural network (RNN) Use the Google stock prices dataset and design a time series analysis and prediction system using RNN.


```
[15]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from tensorflow.keras.preprocessing.sequence import TimeseriesGenerator

[16]: df = pd.read_csv("/Users/sakshisachinpotdar/Downloads/GOOGLE Stock Data set -_
↳GOOGLE Stock Data set.csv")

[17]: df = df[['Date', 'Close']]

[18]: df['Date'] = pd.to_datetime(df['Date'])
df.set_index('Date', inplace=True)

[19]: df['Close'].plot(figsize=(12, 6))
plt.title("Google Stock Price")
plt.ylabel("Price in USD")
plt.show()
```



```
[20]: scaler = MinMaxScaler(feature_range=(0, 1))
scaled_data = scaler.fit_transform(df['Close'].values.reshape(-1, 1))

[21]: look_back = 60 # Use the past 60 days to predict the next day's price
generator = TimeseriesGenerator(scaled_data, scaled_data, length=look_back,
↳batch_size=32)
```

```
[22]: train_size = int(len(df) * 0.8)
      train_data, test_data = scaled_data[:train_size], scaled_data[train_size:]
```

```
[23]: model = Sequential()
      model.add(LSTM(50, activation='relu', input_shape=(look_back, 1)))
      model.add(Dense(1)) # Output layer (single value for next price)
      model.compile(optimizer='adam', loss='mean_squared_error')
```

```
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-
packages/keras/src/layers/rnn/rnn.py:200: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)
```

```
[24]: model.fit(generator, epochs=10, verbose=1)
```

```
# Make predictions
```

Epoch 1/10

```
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-
packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121:
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in
its constructor. `**kwargs` can include `workers`, `use_multiprocessing`,
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be
ignored.
```

```
self._warn_if_super_not_called()
```

```
137/137          1s 6ms/step -
```

```
loss: 0.0839
```

Epoch 2/10

```
137/137          1s 6ms/step -
```

```
loss: 0.0189
```

Epoch 3/10

```
137/137          1s 7ms/step -
```

```
loss: 0.0088
```

Epoch 4/10

```
137/137          1s 7ms/step -
```

```
loss: 9.0410e-04
```

Epoch 5/10

```
137/137          1s 6ms/step -
```

```
loss: 1.9288e-04
```

Epoch 6/10

```
137/137          1s 7ms/step -
```

```
loss: 9.6494e-05
```

Epoch 7/10

```
137/137          1s 7ms/step -
```

```
loss: 8.1387e-05
```

Epoch 8/10

```

137/137          1s 7ms/step -
loss: 1.3209e-04
Epoch 9/10
137/137          1s 7ms/step -
loss: 8.2351e-05
Epoch 10/10
137/137          1s 6ms/step -
loss: 9.2830e-05

```

[24]: <keras.src.callbacks.history.History at 0x159c92810>

```
[25]: predicted_stock_price = model.predict(generator)
```

```

137/137          0s 2ms/step


```

```

[26]: # Inverse transform the predictions back to original scale
predicted_stock_price = scaler.inverse_transform(predicted_stock_price)

# Plot actual vs predicted stock prices
plt.figure(figsize=(12, 6))
plt.plot(df.index[look_back:], scaler.inverse_transform(scaled_data[look_back:
↪]), color='blue', label='Actual Google Stock Price')
plt.plot(df.index[look_back:], predicted_stock_price, color='red', ↪
↪label='Predicted Google Stock Price')
plt.title("Google Stock Price Prediction using LSTM")
plt.xlabel("Date")
plt.ylabel("Stock Price in USD")
plt.legend()
plt.show()

```





[ ]:

[ ]:

[ ]: